

Splay Trees and B-Trees

Namratha M

Assistant Professor,

Department of CSE, BMSCE

Self adjusting Trees





- Ordinary binary search trees have no balance conditions
 - › what you get from insertion order is it
- Balanced trees like AVL trees enforce a balance condition when nodes change
 - › tree is always balanced after an insert or delete
- Self-adjusting trees get reorganized over time as nodes are accessed
 - › Tree adjusts after insert, delete, or find

Splay Trees

- Splay trees are tree structures that:
 - › Are not perfectly balanced all the time
 - › Data most recently accessed is near the root.
(principle of locality; 80-20 “rule”)
- The procedure:
 - › After node X is accessed, perform “splaying” operations to bring X to the root of the tree.
 - › Do this in a way that leaves the tree more balanced as a whole

Splay Tree Terminology


- Let X be a non-root node with ≥ 2 ancestors.
 - P is its parent node.
 - G is its grandparent node.




Zig-Zig and Zig-Zag

Parent and grandparent
in same direction.

Zig-zig




Parent and grandparent
in different directions.



Splay Tree Operations


1. Helpful if nodes contain a **parent** pointer.



2. When X is accessed, apply one of **six** rotation routines.
 - Single Rotations (X has a P (the root) but no G)
ZigFromLeft, ZigFromRight
 - Double Rotations (X has both a P and a G)
ZigZigFromLeft, ZigZigFromRight
ZigZagFromLeft, ZigZagFromRight

Zig at depth 1 (root)


- “Zig” is just a **single rotation**, as in an AVL tree
- Let R be the node that was accessed (e.g. using Find)



- ZigFromLeft moves R to the top → faster access next time

Zig at depth 1


- Suppose Q is now accessed using Find



- `ZigFromRight` moves Q back to the top


Zig-Zag operation

- “Zig-Zag” consists of **two rotations of the opposite direction** (assume R is the node that was accessed)




Zig-Zig operation


- “Zig-Zig” consists of two single rotations of the same direction (R is the node that was accessed)




Decreasing depth - "autobalance"




(a)



(b)



(c)



(d)

Find(T)



Find(R)




Splay Tree Insert and Delete



- Insert x
 - › Insert x as normal then splay x to root.
- Delete x
 - › Splay x to root and remove it. (note: the node does not have to be a leaf or single child node like in BST delete.) Two trees remain, right subtree and left subtree.
 - › Splay the max in the left subtree to the root
 - › Attach the right subtree to the new root of the left subtree.

Example Insert


- Inserting in order 1,2,3,...,8
- Without self-adjustment



With Self-Adjustment

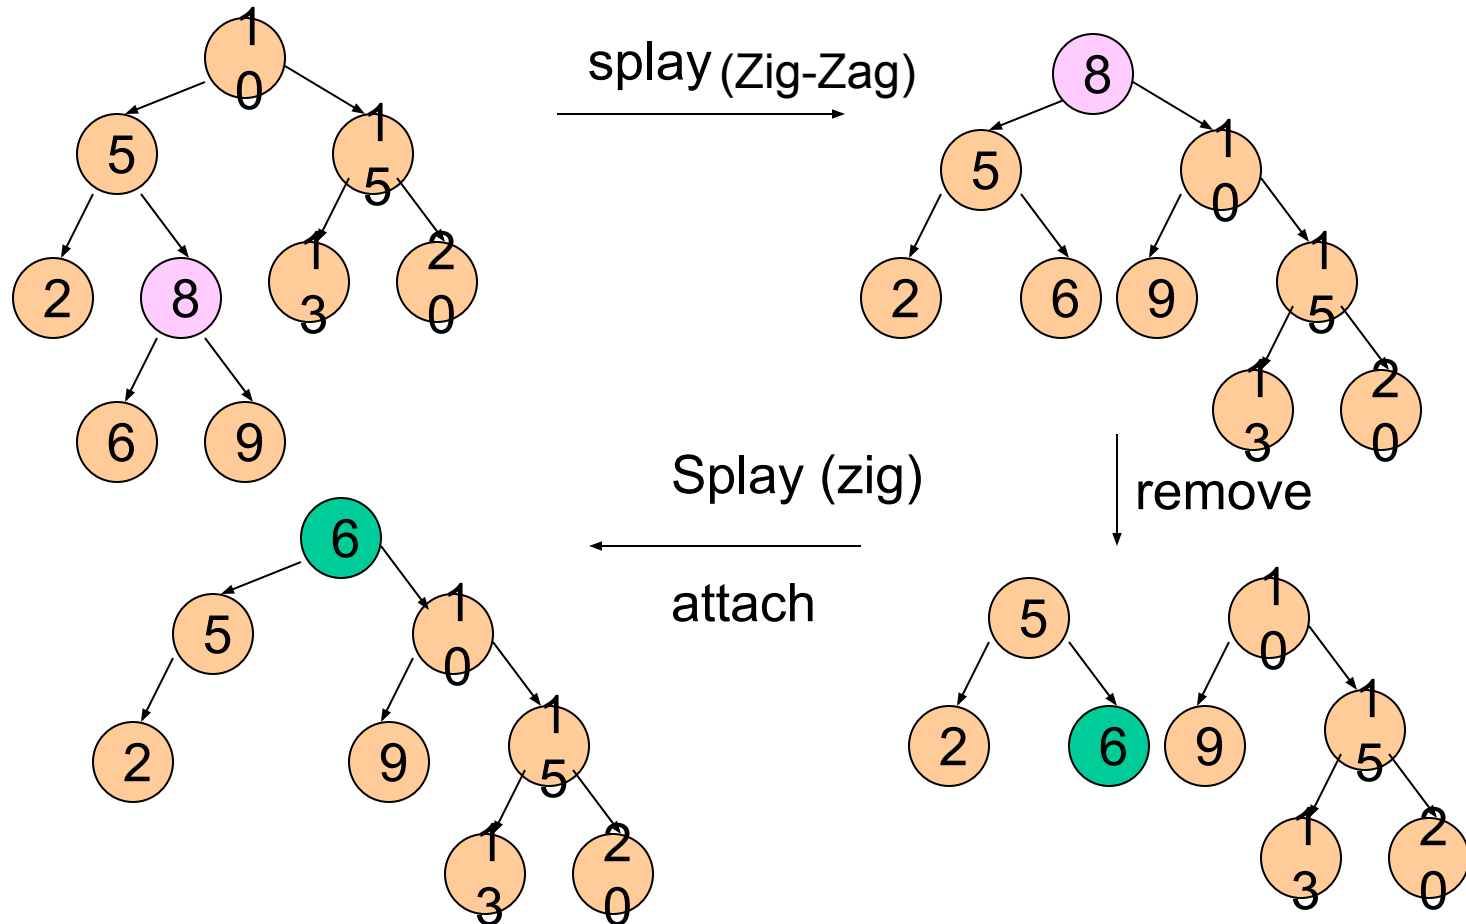


With Self-Adjustment



Each Insert takes $O(1)$ time therefore $O(n)$ time for n Insert!!

Example Deletion




Analysis of Splay Trees

- Splay trees tend to be balanced
 - › M operations takes time $O(M \log N)$ for $M \geq N$ operations on N items. (proof is difficult)
 - › Amortized $O(\log n)$ time.
- Splay trees have good “locality” properties
 - › Recently accessed items are near the root of the tree.
 - › Items near an accessed one are pulled toward the root.

Beyond Binary Search Trees: Multi-Way Trees

- Example: B-tree of order 3 has 2 or 3 children per node



- Search for 8

B-Trees

B-Trees are **multi-way search trees** commonly used in database systems or other applications where data is stored externally on disks and keeping the tree shallow is important.

A **B-Tree of order M** has the following properties:

1. The **root** is either a leaf or has **between 2 and M children**.
2. All nonleaf nodes (except the root) have **between $\lceil M/2 \rceil$ and M children**.
3. **All leaves are at the same depth**.

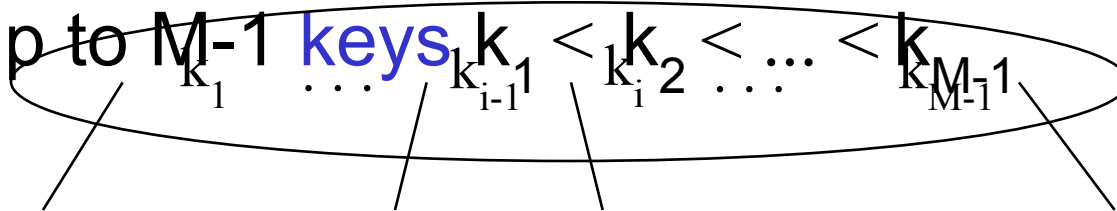
All data records are stored at the leaves.
Internal nodes have “keys” guiding to the leaves.
Leaves store between $\lceil M/2 \rceil$ and M data records.

B-Tree Details

Each (non-leaf) internal node of a B-tree has:

- › Between $\lceil M/2 \rceil$ and M children.

- › up to $M-1$ **keys** $k_1 < k_2 < \dots < k_{M-1}$




The diagram illustrates a B-tree internal node. It contains a sequence of keys: $k_1, \dots, k_{i-1}, k_i, \dots, k_{M-1}$. The word "keys" is highlighted in blue. These keys are enclosed in a horizontal oval. Below the oval, several lines (pointers) extend downwards, representing the children of the node. The lines are positioned such that they fall into the gaps between the keys, indicating that each key separates the pointers to its left and right child.

Keys are ordered so that:

$$k_1 < k_2 < \dots < k_{M-1}$$

Properties of B-Trees



Children of each internal node are "between" the items in that node.

Suppose subtree T_i is the i th child of the node:

all keys in T_i must be between keys k_{i-1} and k_i

i.e. $k_{i-1} \leq T_i < k_i$


k_{i-1} is the smallest key in T_i

All keys in first subtree $T_1 < k_1$

All keys in last subtree $T_M \geq k_{M-1}$

Example: Searching in B-trees


- B-tree of order 3: also known as 2-3 tree (2 to 3 children)



- Examples: Search for 9, 14, 12
- Note: If leaf nodes are connected as a Linked List, B-tree is called a B+ tree – Allows sorted list to be accessed easily


Inserting into B-Trees

- Insert X: Do a Find on X and find appropriate leaf node
 - › If leaf node is not full, fill in empty slot with X
 - E.g. Insert 5
 - › If leaf node is full, **split** leaf node and adjust parents up to root node
 - E.g. Insert 9



Deleting From B-Trees

- Delete X : Do a find and remove from leaf
 - › Leaf underflows – borrow from a neighbor
 - E.g. 11
 - › Leaf underflows and can't borrow – merge nodes, delete parent
 - E.g. 17



Run Time Analysis of B-Tree Operations

- For a B-Tree of order M
 - › Each internal node has up to M-1 keys to search
 - › Each internal node has between $\lceil M/2 \rceil$ and M children
 - › Depth of B-Tree storing N items is $O(\log_{\lceil M/2 \rceil} N)$
- Find: Run time is:
 - › $O(\log M)$ to binary search which branch to take at each node. But M is small compared to N.
 - › Total time to find an item is $O(\text{depth} * \log M) = O(\log N)$

Summary of Search Trees

- Problem with Binary Search Trees: Must keep tree balanced to allow fast access to stored items
- AVL trees: Insert/Delete operations keep tree balanced
- Splay trees: Repeated Find operations produce balanced trees
- Multi-way search trees (e.g. B-Trees): More than two children
 - › per node allows shallow trees; all leaves are at the same depth
 - › keeping tree balanced at all times